

this is around vwap footprint study, apparently hedgefunds look at historical return patterns to predict the returns for real time. Is this possible?

Aye, let's break this down in a clear, realistic, *front-office quant* way — Scottish accent as requested GB 🏴.

Short answer: Yes, it's possible — but only under specific conditions, and only for short-horizon microstructure effects.

It is *not* the same as forecasting mid-term returns like a factor model. It's much closer to *microstructure pattern recognition*.

✓ 1. What hedge funds actually do

Hedge funds (esp. HFT / market-making firms like Citadel, Jane Street, Optiver) *absolutely* use historical return patterns to predict short-horizon returns *in real time*.

But they don't use simple returns — they use **conditional microstructure states**:

✓ Lagged return patterns

e.g.,

- If price drifts up when queue is thin + buy pressure is high,
- Then the probability of further uptick in the next 50–500ms is positive.

✓ "State-dependent" patterns

They don't look at unconditional returns.

They look at patterns *conditional* on:

State variable	Meaning
Depth imbalance	Bid/ask size imbalance
OFI (order flow imbalance)	Market buy vs sell pressure
Queue dynamics	Queue depletion / replenishment
Short-lag autocorrelation	Momentum at 50ms, 100ms, 500ms
Volatility regime	Is market whippy or stable
Spread regime	Wide vs narrow
Event state	Auction, open/close, news

The idea is:

"Given the current state of the order book + recent flow, what are the odds the next

micro-return is up or down?"

This is absolutely feasible and used everywhere: **HFT alpha modelling, market making, signal engines, VWAP/TWAP footprint detection, etc.**

✓ 2. Why this relates to your VWAP footprint study

You're asking:

"If we trade today, is the market reacting the same way as historically?
Can others detect our footprint by matching real-time behaviour to historical patterns?"

This is exactly what HFT firms do:

They continuously compare:

(A) Real-time micro-returns

vs

(B) Historical "normal behaviour" under similar states

vs

(C) Historical "when big flow is present" behaviour

If today's behaviour resembles (C) more than (A), they infer:

"There's someone trading behind the scenes — probably VWAP/POV/IS flow."

This is detection of "flow influence".

✓ 3. Is it theoretically possible?

Yes — because:

✓ Market microstructure has *short-term memory*

Returns over 50ms–2s are **not IID**.

They show:

- Autocorrelation
- Mean reversion
- Order-flow persistence
- Queue depletion patterns

This makes micro-returns *predictable in expectation*, even if only for tiny horizons.

✓ Order flow exhibits persistence

Institutional flow (VWAP/POV) produces a **train of similarly-signed trades**.

This creates detectable statistical signatures:

- Signed volume autocorrelation
- Price impact shape
- Slow reaction of LOB liquidity providers
- Spread dynamics

HFT firms exploit this

They learn patterns like:

- "After a sequence of medium-sized buys at the ask, the next few trades have +X bp drift."
- "If a queue is repeatedly eaten at a constant cadence, there is a VWAP slicing behind it."

They can match your child order signature *to their historical library of flow signatures.*

4. Examples of pattern recognition models used

These are **real models** used on dealer desks, HFT, and hedge funds:

1. ACF of signed returns (micro-momentum)

If:

- r_t = price change
- sign_t = +1 for buy-initiated trades, -1 for sells

Then hedge funds compute:

- $\text{ACF}(\text{sign}_t)$
- $\text{ACF}(r_t)$
- $\text{XCF}(\text{sign}_t, r_t+\text{lag})$
- Autocorrelation of OFI

These show whether incoming flow has predictable direction.

2. Fourier-based pattern matching

This is common for:

- Detecting periodic child order slices
- Identifying constant-cadence VWAP/POV algos

If your slices arrive at a semi-regular interval, the FFT will show peaks at that frequency.

But:

FFT works only when:

- Slicing cadence is stable
- There's low noise
- There's consistent behaviour

VWAP tends to have *variable slicing intervals*, so FFT is sometimes weaker than ACF/XCF.

✖ 3. Machine learning “state-to-return” predictors

Typical models:

- Logistic regression on order-book states
- Random forests on OFI, spread, depth
- LSTMs on event sequences
- Gradient boosted trees (LightGBM / XGBoost)
- HMMs (hidden Markov models) for latent liquidity regimes

Inputs typically:

- Depth imbalance
- Quote replenishment
- Immediate past returns
- Order arrival intensity

HFT desks run these in real-time at 100µs–10ms frequencies.

✓ 5. Bottom line — is it possible to predict returns in real time?

- ✓ Yes — for microstructure horizons (10ms to ~5s).
 - ✓ Yes — especially when flow is strong or systematic.
 - ✓ Yes — by conditioning on order-book states.
- ✗ No — for longer horizons (minutes–hours), markets revert to noisy / semi-random.

VWAP footprints absolutely can be detected like this.

That's why your study is interesting:

You're basically recreating an **HFT flow-detection system** from your algo POV.

✓ 6. If you'd like, I can give you one of the following next:

1. A full statistical pipeline to replicate what HFT firms do
2. KDB+ code to compute state-conditioned returns, ACF, XCF, OFI
3. A detection model to see if your VWAP slices are identifiable
4. FFT approach vs ACF/XCF — detailed comparison and when each works
5. How to construct a “flow signature” library
6. How to test if real-time returns match historical footprint states

Which one do you want, mate?



give me all

Thought for 19s >

Nice, we're going full "build an HFT footprint lab in-house" mode then 😊

I'll go through all 6 pieces you asked for, but keep each one tight and usable.

1 Full statistical pipeline (what a hedge fund would actually do)

Goal:

Given your child orders + full-market L1/L2 + trades, measure whether:

- You can predict short-horizon returns from order-book state, and
- Your VWAP flow leaves a detectable signature others could exploit.

1. Data you need

Per symbol, per venue (ideally):

- **Quotes:** time, sym, bid, bidsz, ask, asksz, maybe depth levels.
- **Trades:** time, sym, px, qty, side (or aggressor flag).
- **Your child orders:** time, sym, side, qty, parent_id, child_id, algo (VWAP/POV/etc).

2. Build microstructure features

On a unified event stream (e.g. trade+quote updates):

For each event:

- Midprice: $\text{mid_t} = (\text{bid} + \text{ask}) / 2$
- Spread: $\text{spr_t} = \text{ask} - \text{bid}$
- Depth imbalance:
 $\text{imb_t} = (\text{bidsz} - \text{asksz}) / (\text{bidsz} + \text{asksz})$
- Order Flow Imbalance (simplified):
 - For trade events:
 - $\text{sign_t} = +1$ if trade at/near ask, -1 if at/near bid
 - $\text{oflow_t} = \text{sign_t} * \text{qty_t}$
 - Volatility proxy: rolling σ_t of mid returns over a short window (e.g. 30s).

3. Define returns you want to predict

Typical short-horizon returns:

- **Event time:**
 $r_{\text{t}}(+k) = (\text{mid}_{\{\text{t}+k\}} - \text{mid}_t) / \text{mid}_t$ for $k = 1, 2, \dots$ events

- **Clock time:**

$$r_t(\Delta) = (\text{mid}_{t+\Delta} - \text{mid}_t) / \text{mid}_t \text{ for } \Delta \in \{100\text{ms}, 1\text{s}, 5\text{s}\}$$

You can look at:

- Plain returns
- Spread-normalised returns: r_t / spr_t
- Vol-normalised returns: r_t / σ_t

4. Baseline vs trade-day (your previous idea formalised)

For each symbol & time-of-day:

- **Baseline:** use N days when you did not trade that symbol
 - Estimate conditional expectations:
 $E[r_t(\Delta) | \text{state}_t = s]$ where state_t is a vector of features (imbalance, OFI, spread, etc.)
- **Trade days:** same symbol + time-of-day, but days you *did* trade.
 - Estimate the same conditional expectation.

Then compare:

Does $E[r_t(\Delta) | \text{state}, \text{trade_day}]$ deviate significantly from $E[r_t(\Delta) | \text{state}, \text{baseline}]$?

If yes, your flow is changing micro-returns relative to "normal".

5. Models to estimate conditional expectations

Start simple:

- **Non-parametric buckets:**
 Bucket by `imbalance` (e.g. [-1,-0.5], ...), `spread`, `vol_regime` and compute average $r_t(\Delta)$.
- **Linear / logistic regression:**
 $r_t(\Delta)$ or $\text{sign}(r_t(\Delta))$ on features: imbalance, OFI, spread, last few returns, etc.
- **Tree models (LightGBM/XGBoost)** for richer interactions.

You can run **separate models** for:

- Baseline periods (no your flow)
- Days with your flow

Then:

- Compare coefficients / feature importances
- Compare predicted vs actual returns
- Test whether residuals behave differently when you trade.

6. Outputs & diagnostics

- **Signal curves:** $E[r_t(\Delta)]$ by OFI / imbalance buckets.
- **ACF/XCF of signed returns / OFI.**
- **Impact shape:** return vs time since your child order.

- “Detectability score”: how well a classifier can distinguish “your-flow-present” vs “no-flow”.
-

2 KDB+/q code skeleton for key microstructure stats

Assume:

- quote : time , sym , bid , bidsz , ask , asksz
- trade : time , sym , px , qty , side ("B" buyer-initiated, "S" seller-initiated)
- child : time , sym , side , qty , parentId , childId

2.1 Basic mid, spread, imbalance from quotes

q Copy code

```
/ Mid, spread, depth imbalance
q)quoteEnriched: update
    mid: (bid + ask) % 2f,
    spr: ask - bid,
    imb: (bidsz - asksz) % (bidsz + asksz)
from quote
```

2.2 Trade sign and signed volume

q Copy code

```
/ +1 for buy-initiated, -1 for sell-initiated
q)tradeEnriched: update
    sign: $[side = "B"; 1; -1],
    signedQty: sign * qty
from trade
```

2.3 Simple OFI (event-time, trade-only version)

q Copy code

```
/ Aggregate trades to, say, 1s buckets per sym
q)tbars: select
    qty: sum qty,
    signedQty: sum signedQty,
    ntrades: count i,
    vwap: sum qty*px % sum qty
by sym, time: 0D00:00:01 xbar time
from tradeEnriched;

/ OFI ~ signed volume per bar
q)tbars: update
    ofi: signedQty
from tbars;
```

You can join in quote features at bar open or close:

q

 Copy code

```
/ Get last quote in each bar
q)qbars: select by sym, time:0D00:00:01 xbar time from quoteEnriched;

/ Join tbars with qbars on sym,time
q)bars: lj[`sym`time; tbars; qbars];
```

Now bars has mid, spr, imb, ofi, etc.

2.4 Returns over horizons

Example: 5s forward mid return:

q

 Copy code

```
/ Assume bars sorted by sym,time
q)bars: `sym`time xasc bars;

/ 5s horizon: use shifted mid per sym
/ number of bars for 5s if bar is 1s = 5
q)fwd5:{[tab;h]
    tab update
    r5: (neg[h] mid - mid) % mid
    by sym
    from tab
}][bars;5];

q)bars: fwd5;
```

2.5 ACF and XCF functions in q (event-time / bar-time)

Say you want ACF of signed returns:

q

 Copy code

```
/ Simple helper: autocorrelation up to maxlag
/ corr is standard correlation in q
q)acf:{[x;maxlag]
    {cor[x[0 + til count x - y]; x[y + til count x - y]]}
    each x enlist til maxlag
};

/ Example per sym on signedQty
q)acfs: select acfVals: enlist acf[signedQty;20] by sym from tbars;
```

Cross-correlation between OFI and returns:

q

 Copy code

```

/ Cross-correlation with lags: x -> y lagged by k
q)xcf:{[x;y;maxlag]
  {cor[x[0 + til count x - y]; y[y + til count x - y]]}
  each x enlist til maxlag
};

/ Example: ofi vs r5
q)xcfVals: select xcfVals: enlist xcf[ofi; r5; 20] by sym from bars;

```

2.6 Join your child orders and compute impact

q Copy code

```

/ Nearest mid at child send time (previous quote)
q)childWithMid: lj[`sym`time;
  child;
  select last mid by sym, time: 0D00:00:01 xbar time from quoteEnriched
];

/ 5s post mid for each child: via as-of join to bars with r5
q)childImpact: aj[`sym`time; childWithMid; select sym,time,r5 from bars];

/ Now childImpact has r5 = 5s mid return post child slice

```

You can then bucket r5 by spread/OFI/imbalance.

3 Detection model: "Is there VWAP-like flow present?"

Think of this as a **classification** problem.

3.1 Labels

Define per bar (or per second):

- $y_t = 1$ if your algo is active in that symbol at that time (you have child orders overlapping that bar)
- $y_t = 0$ otherwise (baseline periods)

3.2 Features

Per (sym, time bar):

- OFI, depth imbalance
- Spread, mid-level
- Recent returns (e.g. r1, r5, r10)
- Trade intensity (trades per second, volume)
- Buy/sell imbalance (fraction of buy-initiated trades)
- Quote update intensity
- Your parent direction, if known (for your own internal tests)

3.3 Model choices

Start very simple:

- Logistic regression:
 - $P(\text{flow_active}=1 \mid \text{features})$
- Gradient boosting (LightGBM/XGBoost) for non-linearities
- Optional: HMM/regime models if you want “algo active vs not” as latent.

3.4 Evaluation

- Train on historical data where you *know* when you traded.
- Check:
 - ROC-AUC for predicting y_t .
 - Precision / recall at reasonable thresholds.
 - Whether model generalises to new days.

If an external HFT had the same data, and your model can easily separate $y=1$ vs 0 , then **your flow is detectable**.

4 FFT vs ACF/XCF – when and why

You'd asked this explicitly before; cleaned-up answer:

ACF / XCF (time-domain)

- ACF: correlation of a signal with itself at different lags.
 - Good for:
 - Detecting **serial dependence** (momentum / mean reversion).
 - “How persistent is signed volume, returns, OFI?”
- XCF: correlation between two signals at different lags.
 - Good for:
 - “Does OFI predict returns 1–10s ahead?”
 - “Do our child order timestamps align with spikes in OFI/returns?”

Pros:

- Directly linked to economic intuition.
- Easy to interpret (e.g. positive at lag +5 → our flow affects returns after 5 bars).
- Works well with irregular and noisy flows.

Cons:

- You must pick max lags.
- Doesn't highlight *strict* periodicities as crisply as FFT.

FFT (frequency-domain)

FFT decomposes your time series into sinusoidal components.

- Good for:
 - Detecting **periodicity** / regular slicing cadence.

- e.g. child order every 2s → strong spectral peak at 0.5 Hz.

Pros:

- Great when the algo slices with quasi-constant cadence.
- Very clear for “metronomic” TWAP, some POVs, some DMA clients.

Cons:

- Sensitive to irregular timing and noise.
- VWAP often has **variable slice size and timing** → spectral lines blur.
- Harder to link to economic microstructure interpretation.

Which is more practical for your VWAP footprint work?

- For impact / predictability of returns → ACF/XCF is **more practical**.
- For detecting a fixed slicing frequency pattern → **FFT can help**, but only if your flow is quite regular.

If you must prioritise:

1. Start with ACF/XCF of:
 - trade sign
 - OFI
 - returns
 - your child order indicator
2. Add FFT later if you explicitly suspect quasi-periodic slicing.

5 Flow “signature library” – how to build it

Think like an HFT:

For each *type of flow* (VWAP buy, VWAP sell, POV, IS, etc.), what does the local market typically look like before / during / after?

5.1 Define signature objects

For each (algo type, side, symbol/sector, volatility regime), build:

- **Impact curve**: average $r(\Delta)$ for $\Delta \in \{1s, 5s, 30s, 300s\}$.
- **Spread response**: average change in spread around your child orders.
- **Depth response**: how L1/L2 depth changes before vs after.
- **OFI profile**: signed volume pattern around your flow.
- **Autocorrelation profile**: ACF of signed trades around parent order.

Store these as templates:

text

 Copy code

```
signature[algo="VWAP", side="Buy", volRegime="High", symBucket="LargeCap"]
  .impactCurve[Δ]
```

```
.spreadCurve[Δ]
.ofiCurve[Δ]
.acfSign[lag]
.acfReturn[lag]
```

5.2 How to build them

For every historical parent order:

1. Align time to $t=0$ at first child slice.
2. For $Δ ∈ [-T_{pre}, …, T_{post}]$:
 - Record features: mid, spread, OFI, depth, etc.
3. Normalise:
 - Express returns in bps or spread units.
 - Maybe normalise by volatility regime.
4. Average across many parents (within each bucket).

This gives you **mean curves + confidence bands**.

6 Real-time matching of current behaviour to historical signatures

Now the fun bit: "Can someone detect us in real time?"

6.1 Build a live feature window

At each timestamp t (e.g. every second):

- Build a **feature vector** from recent history $[t-T, t]$:
 - Sequence of OFI, mid returns, spread, depth imbalance.
 - Maybe compressed into summary stats:
 - cumulative OFI
 - realised volatility
 - skewness of returns
 - fraction of buy-initiated trades
 - trend in mid over last N seconds

6.2 Compare against signature templates

Two broad approaches:

(a) Distance / similarity

For each template (VWAP-buy, VWAP-sell, etc.):

- Compute distance between:
 - **Current feature profile** and
 - **Template mean curve**,

using e.g.:

- Euclidean distance
- Cosine similarity
- Dynamic Time Warping (DTW) if timing is a bit misaligned.

If distance is small and persistent over several seconds, label:

"Likely VWAP-like buying flow present."

You can do this internal to see how "obvious" your flow looks.

(b) Supervised classifier (more HFT-like)

You already defined labels (`flow_active` vs `baseline`).

- Train model on historical data to predict `flow_active`.
- In real time, feed in the live feature vector → get $P(\text{flow_active}=1 \mid \text{state}_t)$.

If this probability is high and stable for a bit, then:

- A hypothetical external HFT with similar data could identify your flow.

How this all ties back to VWAP footprint & "is it possible?"

Putting it together:

1. Yes, it's possible to use historical patterns to predict short-horizon returns and detect flow.
2. The pipeline above is essentially how:
 - You learn from historical microstructure behaviour (baseline vs your-flow),
 - Build signatures, and
 - Run real-time state-to-return + flow-detection models.
3. If your own internal models can easily spot your flow / predict micro-returns due to it, then external HFTs with comparable data can too.

If you want, next step I can:

- Turn this into a **concrete notebook plan** (Python + PyKX) with:
 - precise table schemas,
 - step-by-step feature engineering,
 - ACF/XCF plots,
 - a simple logistic model for detection;
- or a **kdb-only "lab script"** you can run on a subset of names to prototype quickly.



yes

Cool, let's actually *build the lab* now rather than just talk about it 😊

I'll give you **both**:

1. A **Python notebook plan** (with code) – pandas + sklearn, optionally PyKX entry point.
2. A **kdb+/q "lab"** script to run directly on your tick and child-order tables.

You can then wire to your own infra.

1 Python notebook: VWAP footprint & flow detection lab

1. Imports & basic setup

python Copy code

```
# Core
import numpy as np
import pandas as pd

# Stats / modelling
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score, precision_recall_curve
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline

# Plotting
import matplotlib.pyplot as plt

# Optional: statsmodels for ACF
from statsmodels.tsa.stattools import acf

# Optional: PyKX if you pull from kdb directly
# import pykx as kx
```

2. Data model assumptions

Assume you can bring data into pandas in roughly these shapes:

python Copy code

```
# Quotes: one row per quote update
# time is ns or ms resolution; convert to pandas.Timestamp
quotes = pd.DataFrame({
    # 'time': ...,
    # 'sym': ...,
    # 'bid': ...,
    # 'bidsz': ...,
    # 'ask': ...,
    # 'asksz': ...
})

# Trades: one row per trade
trades = pd.DataFrame({
    # 'time': ...,
    # 'sym': ...,
    # 'px': ...
})
```

```

# 'qty': ...,
# 'side': ... # 'B' buy initiated, 'S' sell initiated
})

# Your child orders
child = pd.DataFrame({
    # 'time': ...,
    # 'sym': ...,
    # 'side': ... # 'B' or 'S' for parent direction or child direction
    # 'qty': ...,
    # 'parent_id': ...,
    # 'child_id': ...,
    # 'parent_algo': ... # 'VWAP', 'POV', etc (optional but nice)
})

```

Convert time to index / datetime:

python

 Copy code

```

for df in [quotes, trades, child]:
    df['time'] = pd.to_datetime(df['time'])

```

Filter to one symbol to prototype:

python

 Copy code

```

sym = 'AAPL' # example
q_sym = quotes[quotes['sym'] == sym].copy()
t_sym = trades[trades['sym'] == sym].copy()
c_sym = child[child['sym'] == sym].copy()

```

3. Enrich quotes: mid, spread, imbalance

python

 Copy code

```

q_sym['mid'] = (q_sym['bid'] + q_sym['ask']) / 2.0
q_sym['spr'] = q_sym['ask'] - q_sym['bid']
q_sym['imb'] = (q_sym['bidsz'] - q_sym['asksz']) / (q_sym['bidsz'] + q_sym['asksz'])

```



4. Enrich trades: sign, signed volume

python

 Copy code

```

t_sym['sign'] = np.where(t_sym['side'] == 'B', 1, -1)
t_sym['signed_qty'] = t_sym['sign'] * t_sym['qty']

```

5. Build 1-second bars combining trades & quotes

You can do this many ways; here's a straightforward approach:

python

 Copy code

```
# Set index for resampling
q_sym = q_sym.set_index('time').sort_index()
t_sym = t_sym.set_index('time').sort_index()

# Quote snapshot per second: take *last* quote in the bucket
q_bar = q_sym.resample('1S').last()[['mid', 'spr', 'imb']]

# Trade aggregates per second
def vwap(px, qty):
    return (px * qty).sum() / qty.sum() if qty.sum() > 0 else np.nan

agg = {
    'px': vwap,
    'qty': 'sum',
    'signed_qty': 'sum',
}

t_bar = t_sym.resample('1S').apply(agg)
t_bar.rename(columns={'px': 'vwap', 'qty': 'vol', 'signed_qty': 'ofi'}, inplace=True

# Combine into a single DataFrame
bars = q_bar.join(t_bar, how='outer')
bars['sym'] = sym

# Fill forward quotes where missing; leave trades as is
bars[['mid', 'spr', 'imb']] = bars[['mid', 'spr', 'imb']].ffill()
bars['vol'] = bars['vol'].fillna(0.0)
bars['ofi'] = bars['ofi'].fillna(0.0)
```

Now bars has:

- mid, spr, imb, vwap, vol, ofi for each second.

6. Compute short-horizon returns

Example: r_1s, r_5s, r_30s forward mid returns.

python

 Copy code

```
for horizon, label in [(1, 'r_1s'), (5, 'r_5s'), (30, 'r_30s')]:
    fwd_mid = bars['mid'].shift(-horizon)
    bars[label] = (fwd_mid - bars['mid']) / bars['mid']
    # Normalise by spread if you like:
    bars[f'{label}_spr'] = bars[label] / bars['spr']
```

7. Label bars as “flow active” vs baseline

We map child orders to 1-second bars and set a binary label.

python

 Copy code

```
# Child orders to 1S buckets
c_sym = c_sym.set_index('time').sort_index()
c_bar = c_sym.resample('1S').agg({'qty': 'sum'})

# Label whether we were active in that second
bars['flow_active'] = (c_bar['qty'] > 0).reindex(bars.index, fill_value=False)

# Optional: also keep parent direction
# If you want direction = sign of net parent/child qty that second:
bars['flow_dir'] = np.sign(c_bar['qty']).reindex(bars.index)
```

8. Split baseline vs trade-day populations

You can do this by days where you trade vs days you don't, or by `flow_active`.

Simple version:

python

 Copy code

```
baseline = bars[~bars['flow_active']].copy()
trade_period = bars[bars['flow_active']].copy()
```

For your original *study* concept, you'd probably:

- Identify *days* with no flow in that symbol → baseline days.
- Identify days with flow → treatment days.
- Tag each bar with `day_has_flow` flag and use that as the grouping.

9. ACF/XCF analysis

9.1 ACF of returns or OFI

python

 Copy code

```
# Drop NaNs
ret_series = bars['r_1s'].dropna()
ofi_series = bars['ofi'].loc[ret_series.index]

max_lag = 20
acf_ret = acf(ret_series, nlags=max_lag, fft=True)
acf_ofi = acf(ofi_series, nlags=max_lag, fft=True)
```

Plot:

```
python Copy code

lags = np.arange(max_lag + 1)

plt.figure()
plt.stem(lags, acf_ret, use_line_collection=True)
plt.xlabel('Lag (seconds)')
plt.ylabel('ACF of r_1s')
plt.title(f'ACF of 1s returns - {sym}')
plt.show()
```

9.2 XCF (OFI → returns)

No built-in, but you can roll your own:

```
python Copy code

def xcf(x, y, max_lag):
    x = np.asarray(x)
    y = np.asarray(y)
    n = len(x)
    x = (x - x.mean()) / x.std()
    y = (y - y.mean()) / y.std()
    corrs = []
    for lag in range(max_lag + 1):
        corrs.append(np.mean(x[:n-lag] * y[lag:]))
    return np.array(corrs)

ret = ret_series.values
ofi = ofi_series.values

xcf_ofi_ret = xcf(ofi, ret, max_lag=20)

plt.figure()
plt.stem(np.arange(21), xcf_ofi_ret, use_line_collection=True)
plt.xlabel('Lag (seconds)')
plt.ylabel('XCF: OFI → r_1s')
plt.title(f'Cross-correlation OFI to 1s returns - {sym}')
plt.show()
```

Do this for **baseline** vs **trade-period** and compare curves.

10. Simple flow detection classifier (hedge-fund style)

Predict `flow_active` from state variables. If this works well, your flow is detectable.

```
python Copy code
```

```
# Feature matrix X, target y
feature_cols = [
    'spr',
    'imb',
    'ofi',
    'vol',
    'r_1s', # previous returns - use Lagged versions
]

# You should lag returns so you're not using forward info:
bars['r_1s_lag1'] = bars['r_1s'].shift(1)
feature_cols = ['spr', 'imb', 'ofi', 'vol', 'r_1s_lag1']

df_model = bars.dropna(subset=feature_cols + ['flow_active'])
X = df_model[feature_cols].values
y = df_model['flow_active'].astype(int).values

X_train = X # for prototype, you can do a true train/test split by day later
y_train = y
```

Model:

python

 Copy code

```
clf = Pipeline([
    ('scaler', StandardScaler()),
    ('logit', LogisticRegression(max_iter=1000))
])

clf.fit(X_train, y_train)
y_score = clf.predict_proba(X_train)[:, 1]

auc = roc_auc_score(y_train, y_score)
print(f"ROC-AUC for detecting your flow: {auc:.3f}")
```

If $AUC \gg 0.5$ (e.g. 0.7–0.9), then **given your features**, your flow is statistically detectable.

11. Distance / similarity to historical “flow signatures”

Suppose you've built **average OFI + return profiles** around your flow (precomputed templates). For now, just show how to compute a rolling feature vector and compare with a template.

11.1 Build rolling feature window

For each time t , for the last T seconds, stack OFI & returns into a vector:

python

 Copy code

```
window = 10 # seconds
```

```
# Example narrow features: last 10s OFI and returns
ofi_seq = np.stack([
    bars['ofi'].shift(lag)
    for lag in range(window)
], axis=1)

r1_seq = np.stack([
    bars['r_1s'].shift(lag)
    for lag in range(window)
], axis=1)

feature_seq = np.hstack([ofi_seq, r1_seq])
feature_seq = pd.DataFrame(
    feature_seq,
    index=bars.index,
    columns=[f'ofi_t-{k}' for k in range(window)] + [f'r1_t-{k}' for k in range(w
)
)
```

11.2 Create a template vector from known VWAP intervals

Say you know intervals where you traded with a VWAP parent; you can average their sequences:

python

 Copy code

```
# Use only times when flow_active is True to define your signature
sig_df = feature_seq[bars['flow_active']].dropna()

template_vec = sig_df.mean(axis=0).values # 1D vector
template_norm = template_vec / np.linalg.norm(template_vec)
```

11.3 Compute cosine similarity in real-time

python

 Copy code

```
# Cosine similarity of each window to the template
seq = feature_seq.dropna()
seq_arr = seq.values
seq_norm = seq_arr / np.linalg.norm(seq_arr, axis=1, keepdims=True)

cos_sim = (seq_norm @ template_norm)

cos_sim_series = pd.Series(cos_sim, index=seq.index)
```

You can now:

- Plot `cos_sim_series` over time.
- Flag intervals where `cos_sim > threshold` (e.g. 0.8) as "VWAP-like flow signature".

2 kdb+/q "VWAP footprint lab" script

Below is a q script you can adapt. It assumes:

- quote and trade tables exist for some date range.
- child table exists with your child orders.

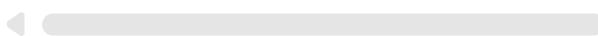
You can run it symbol by symbol or across a universe.

2.1 Basic table schema assumptions

q

 Copy code

```
/ Quotes
/ schema: quote: ([] time:`timestamp$(); sym:`symbol$(); bid:`float$(); bidsz:`float$());
/ trades
/ trade: ([] time:`timestamp$(); sym:`symbol$(); px:`float$(); qty:`float$(); side:`string$());
/ Child orders
/ child: ([] time:`timestamp$(); sym:`symbol$(); side:`symbol$(); qty:`float$(); price:`float$());
```



2.2 Enrich quotes

q

 Copy code

```
/ mid, spread, imbalance
quoteEnriched: update
    mid:(bid + ask) % 2f,
    spr:ask - bid,
    imb:(bidsz - asksz) % (bidsz + asksz)
from quote;
```

2.3 Enrich trades: sign & signed volume

q

 Copy code

```
/ sign: +1 for buy-initiated, -1 for sell-initiated
tradeEnriched: update
    sign:$[side=`B; 1f; -1f],
    signedQty: sign * qty
from trade;
```

2.4 1-second bars from trades

q

 Copy code

```
/ Helper VWAP aggregate
vwap:{[px;qty] $[0f=sum qty; (px * qty) wsum qty; 0n]}

/ Aggregate trades to 1-second buckets
tbars: select
    vwap:vwap[px;qty],
    vol:sum qty,
    ofi:sum signedQty,
    ntrades:count i
    by sym, time:0D00:00:01 xbar time
    from tradeEnriched;
```

2.5 Quote per-second snapshots

q

Copy code

```
/ Last quote in each 1-second bucket
qbars: select last mid, last spr, last imb
    by sym, time:0D00:00:01 xbar time
    from quoteEnriched;
```

2.6 Combine quote + trade bars

q

Copy code

```
/ Left join trades onto quotes or vice-versa - I'll use full outer via uj
/ But we'll keep one canonical table 'bars'.

bars: uj[
    select sym,time,mid,spr,imb from qbars;
    select sym,time,vwap,vol,ofi,ntrades from tbars
];
    
```

/ Sort for lagging

bars: `sym`time xasc bars;

/ Fill forward quotes, default 0 for vol/ofi/ntrades

bars: update
 mid: mid ffill,
 spr: spr ffill,
 imb: imb ffill,
 vol: 0f ^ vol,
 ofi: 0f ^ ofi,
 ntrades: 0 ^ ntrades
 from bars;

2.7 Forward returns in q

Example: 1s, 5s, 30s forward returns by sym.

q

□ Copy code

```
/ Helper: forward horizons in bars (assuming 1-second bars)
/ h is integer number of bars
fwdRet:{[t;h]
    update r:(neg[h] mid - mid) % mid
    by sym
    from t
};

/ Compute forward returns in place
bars: update r1s:(neg[1] mid - mid) % mid by sym from bars;
bars: update r5s:(neg[5] mid - mid) % mid by sym from bars;
bars: update r30s:(neg[30] mid - mid) % mid by sym from bars;

/ Optional: spread-normalised
bars: update r1s_spr: r1s % spr from bars;
```

2.8 Label bars with your flow activity

Map child orders to the 1-second grid.

q

□ Copy code

```
/ Aggregate child orders into 1-second buckets
childBars: select
    childVol: sum qty,
    childDir: sum ${side=`B; qty; -qty] / net directional volume
by sym, time:0D00:00:01 xbar time
from child;

/ Join into bars
bars: lj[`sym`time; bars; childBars];

/ Flag flow active
bars: update
    flowActive: 0f < childVol,           / bool
    flowDir: ${[childVol=0f; 0f; sign childDir] / +1 buy, -1 sell, 0 none
from bars;
```

2.9 ACF/XCF functions in q

q

□ Copy code

```

/ Autocorrelation up to maxlag (returns list of length maxlag+1)
acf: {[x;maxlag]
  x: x where not null x;
  n: count x;
  m: avg x;
  xs: x - m;
  denom: sum xs*xs;
  / lag 0..maxlag
  { (sum xs[til n-y]*xs[y+til n-y]) % denom } each til maxlag+1
};

/ Cross-correlation x -> y
xcf: {[x;y;maxlag]
  / align & drop nulls
  x: x where not null x;
  y: y where not null y;
  n: min[count each (enlist x; enlist y)];
  x: x til n;
  y: y til n;
  mx: avg x;
  my: avg y;
  xs: x - mx;
  ys: y - my;
  denom: sqrt sum xs*xs * sum ys*ys;
  { (sum xs[til n-y]*ys[y+til n-y]) % denom } each til maxlag+1
};

```

Use it:

q Copy code

```

/ Example per sym ACF of 1s returns, OFI, etc.
maxlag: 20;

retAcf: select acfVals: enlist acf[r1s; maxlag] by sym from bars;
ofiAcf: select acfVals: enlist acf[ofi; maxlag] by sym from bars;

/ Example XCF: OFI -> r1s
ofiR1xcf: select xcfVals: enlist xcf[ofi; r1s; maxlag] by sym from bars;

```

2.10 Impact around child orders (VWAP footprint style)

Compute average return after your child slices.

q Copy code

```

/ Function: impact around child orders at given horizon h (bars)
impactAt: {[bars;h]
  / pick timestamps where flowActive just turned on (or where childVol>0)

```

```

co: select from bars where flowActive & 0=prev flowActive by sym;
/ for each such event, look up forward r1s or r5s etc.
/ Using r5s as example
rname:`$"r",string h,"s"; / e.g. `r5s
select avgRet: avg .[bars; (;rname); ()] by sym from
    select sym,time,ret:.[bars; `sym`time!(sym; time); rname] from co
};

/ Example: 5s impact
impact5: impactAt[bars;5];

```

If you want a full curve ($\Delta = 1, 2, 5, 10, 30$ seconds) you can:

- Precompute r1s, r2s, r5s, r10s, r30s
 - Aggregate avg per horizon & per side.
-

2.11 Simple “detectability” score in q (bucket-level)

Instead of full logistic regression in q, you can:

1. Bucket features (e.g. OFI and imbalance).
2. For each bucket, compute:
 - Fraction of time flow is active.
 - Conditional returns.

If the empirical $P(\text{flowActive}=1 \mid \text{bucket})$ varies a lot across buckets, your flow is conditionally predictable.

q

 Copy code

```

/ Example: bucket OFI and imbalance
bucket:{[x;edges] edges bin x}

/ define edges
ofiEdges:-1 -0.5 -0.1 0 0.1 0.5 1;
imbEdges:-1 -0.5 -0.1 0 0.1 0.5 1;

bars: update
    ofiBucket: bucket[ofi; ofiEdges],
    imbBucket: bucket[imb; imbEdges]
from bars;

/ Group stats
detectStats: select
    pFlow: avg flowActive,           / probability flow active
    avgRet1s: avg r1s,
    avgOfi: avg ofi
by sym, ofiBucket, imbBucket
from bars;

```

/ You can inspect detectStats buckets where pFlow is high vs low.

If there are very distinct “high pFlow” buckets that line up with certain imbalance/OFI states, then someone who knows that mapping can infer your presence.

What you have now

You now have:

- **Python notebook skeleton** for:
 - Bar construction
 - Feature engineering
 - ACF/XCF
 - Flow detection via logistic regression
 - Cosine-similarity “signature” matching
- **kdb+/q lab script** to:
 - Build 1s bars
 - Compute mid/spread/imbalance/OFI
 - Compute forward returns